The background of the image features a vibrant, abstract sunset or sunrise. The left side is dominated by warm, fiery orange and yellow tones, while the right side transitions into cooler, ethereal blues and purples. The colors are blended together in a soft, hazy manner, creating a dreamlike atmosphere.

Modular Monolith

Microservices



Microservices

An **architectural style**
that structures an **application**
as a collection of **small,**
loosely coupled, autonomous services,
organized around **busine\$\$ capabilities,**
owned by a small team

History

- **Early 2000s:** Amazon and Netflix started breaking down their monolithic apps into smaller more manageable pieces
- **2011-2012:** Martin Fowler and James Lewis started talking about "Microservices"
- **2014:** M.Fowler and J.Lewis popularized it via talks and terms
- **2015-2020:** Organizations started adopting them + Docker/k8s
- **2018-...:** Microservices is a hype-driven decision !
- **Today:** mainstream; serverless and service mesh appeared

Benefits of Microservices



- ✓ **Faster Time-to-Market** => 😊 Business**
if independently deployable by autonomous👑 teams
- ✓ **Lower Cognitive Load** => 😊 Developers**
if small, business-aligned microservices with clear boundaries
- ✓ **Scalability** for the hot🔥 parts that require it
- ✓ **Availability**: fault-tolerance to partial failures
- ✓ **Technology Stack Freedom**
Safer to upgrade/change the language/library
- ✓ **Security / Privacy (GDPR) / Compliance**

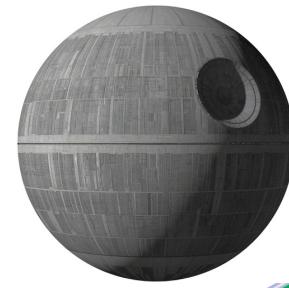
** a Modular Monolith can also give you these

Drawbacks of Microservices

- Development Complexity
- Eventual Consistency
- Network Latency & Reliability
- Deployment Rampage
- Distributed Tracing & Debugging

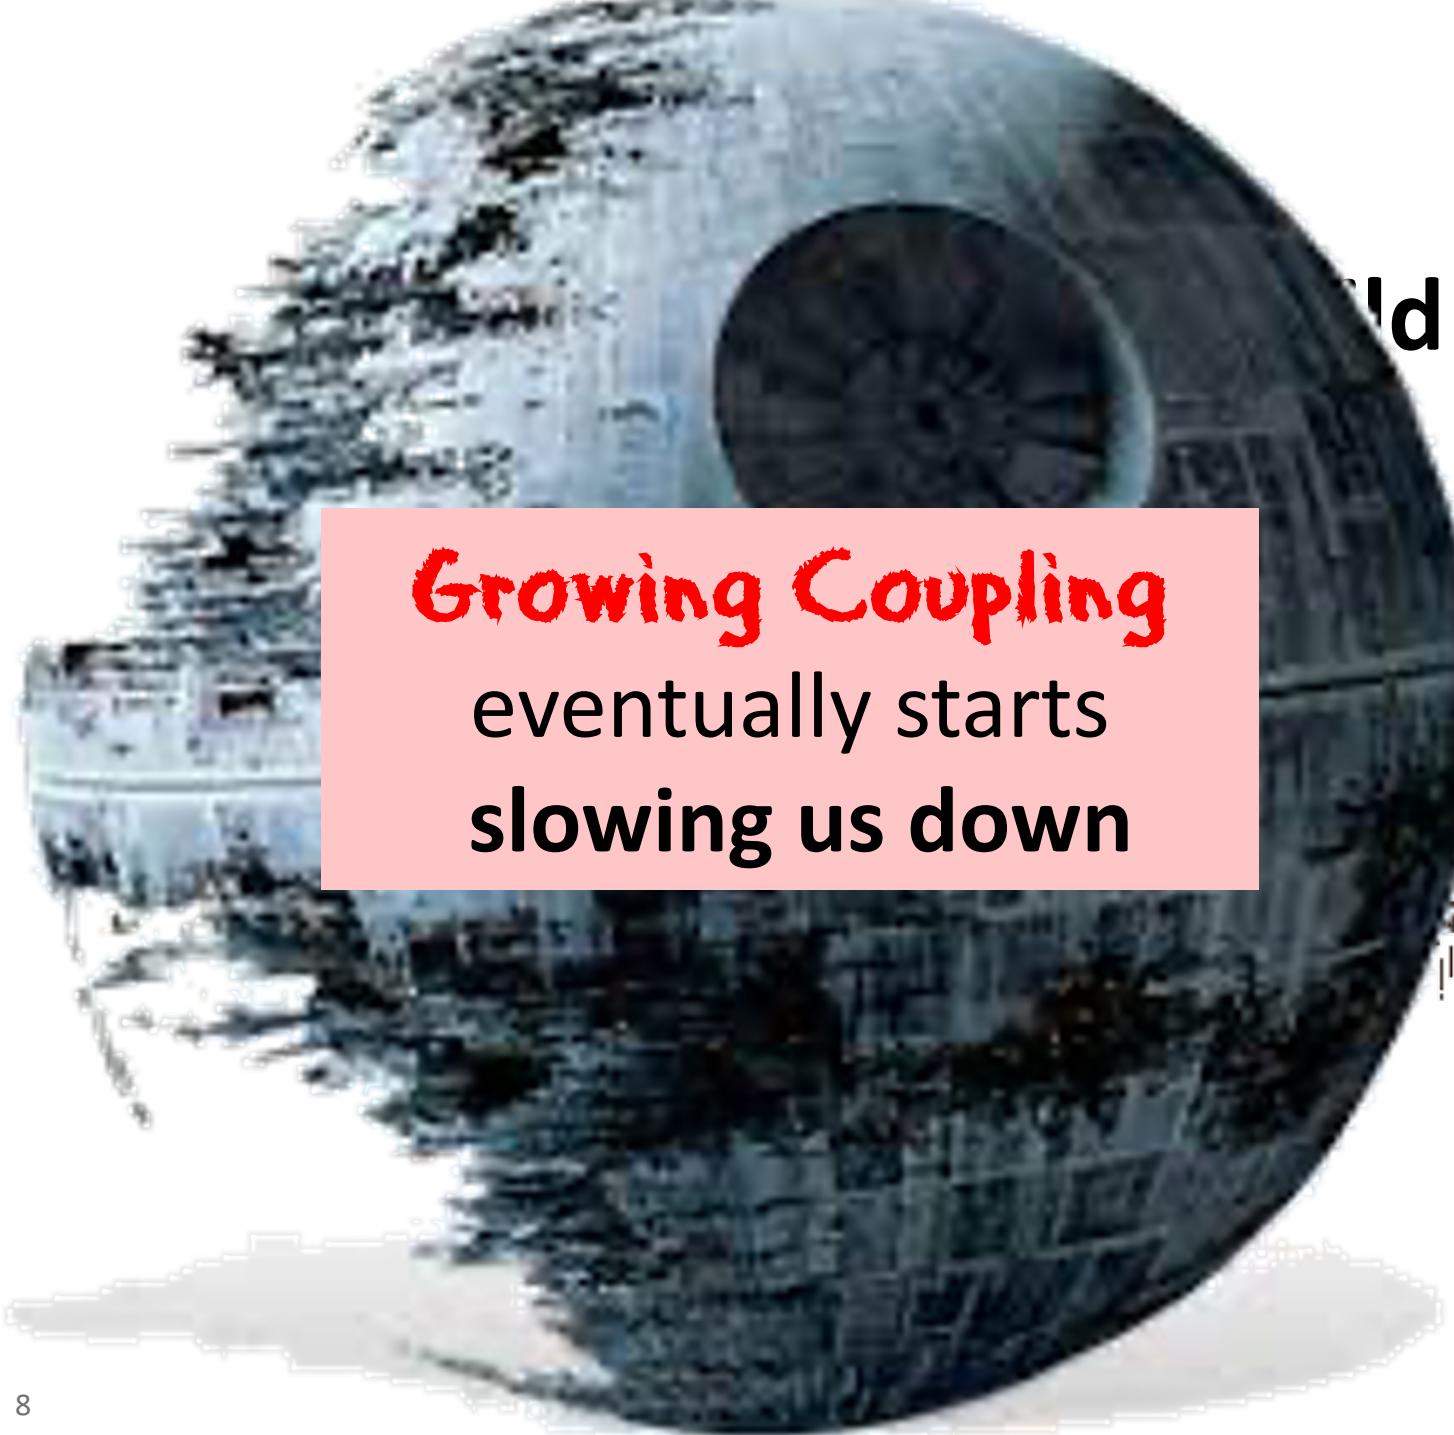
What's faster to build from scratch?

One Monolith



or 3 Microservices





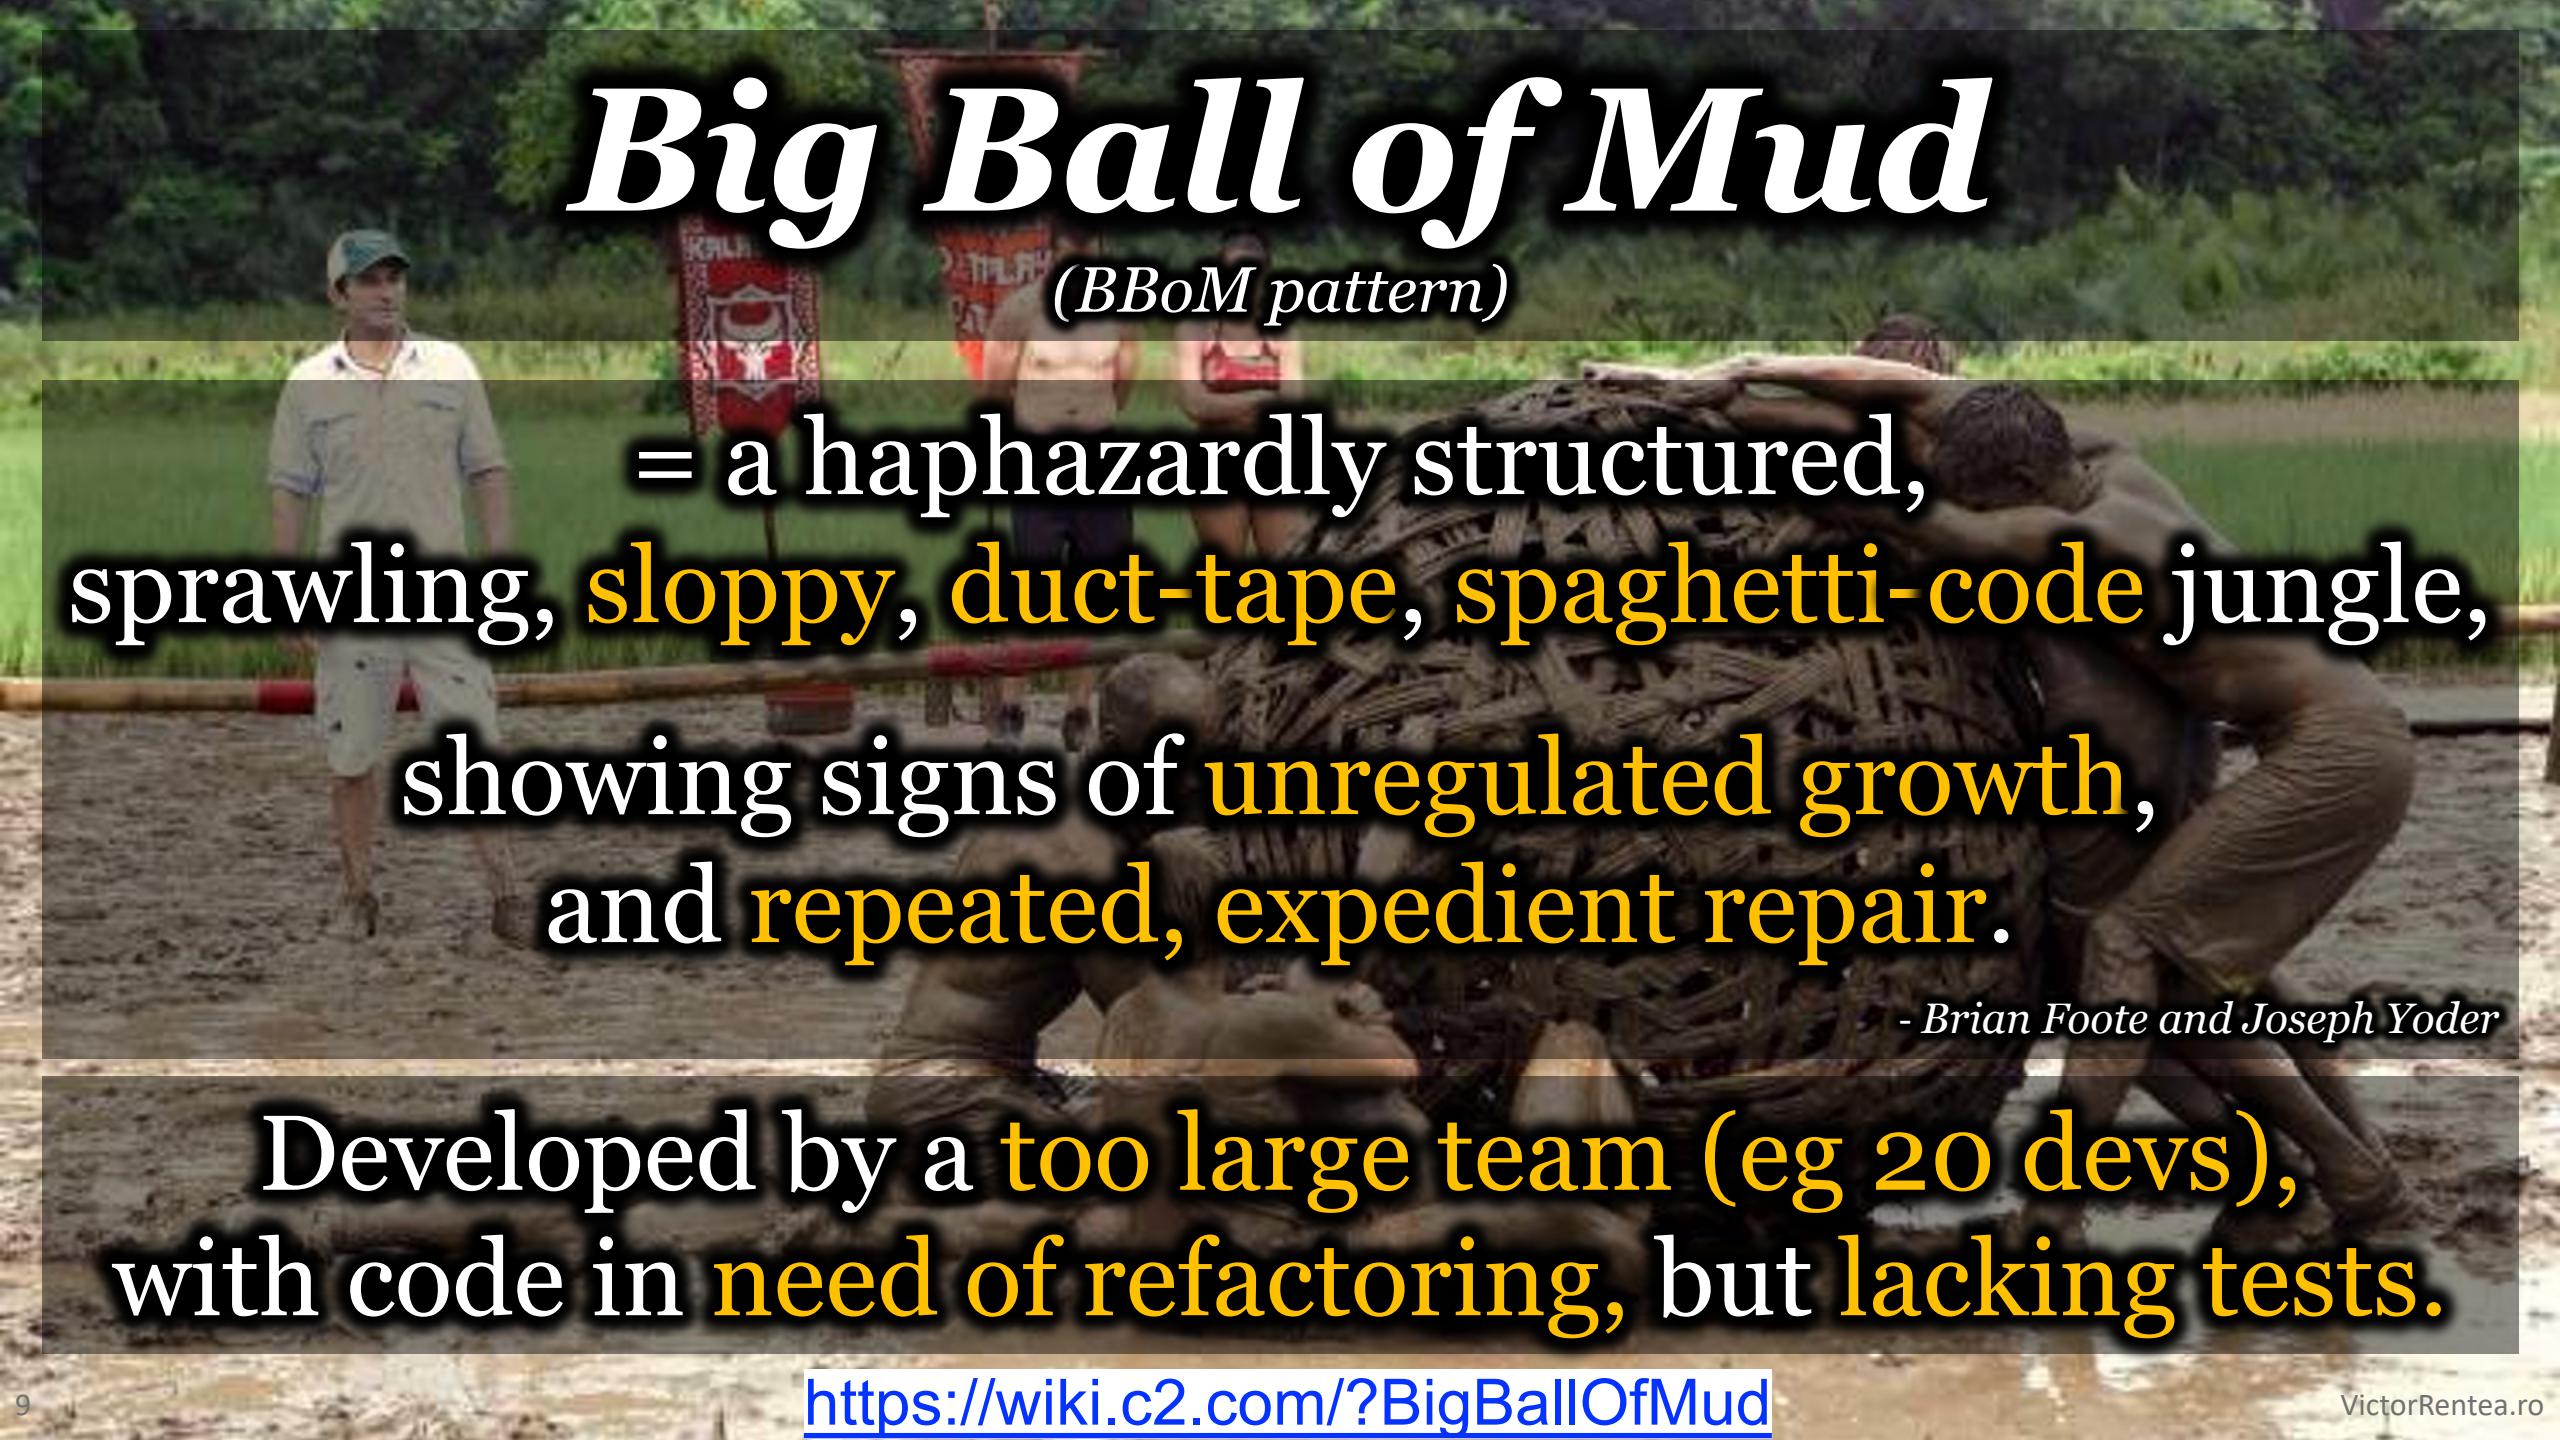
Would from scratch?

Microservices



Big Ball of Mud

(BBoM pattern)



= a haphazardly structured,
sprawling, sloppy, **duct-tape**, spaghetti-code jungle,
showing signs of **unregulated growth**,
and **repeated, expedient repair**.

- Brian Foote and Joseph Yoder

Developed by a too large team (eg 20 devs),
with code in need of refactoring, but lacking tests.

From a BBoM to Microservices - Strategies

1) Big-Bang Rewrite of a 2M lines, 12 years-old codebase

- Start from scratch, re-gather all requirements
- Changes to old system: (a) reject, (b) delay 1 year, or (b) cost*2 (old+new)
- **20-30% success rate. Worse: likely to turn into a Distributed Monolith** 😱



2) Strangler-Fig Pattern (**outside-in**) ✓

- Move behind a proxy **simple features** to separate microservices



3) Split in Modules > Extract (**inside-out**) ✓

- **Progressively decouple** logic and data in the monolithic codebase

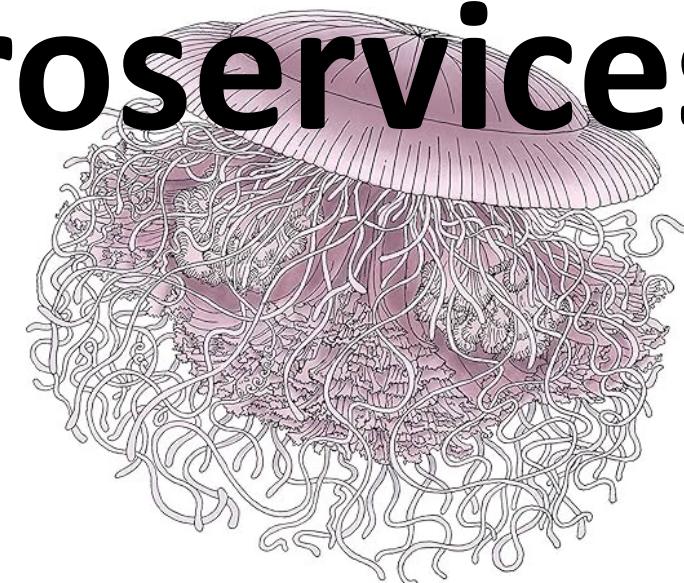


Monolith to Microservices

O'REILLY®

Monolith to Microservices

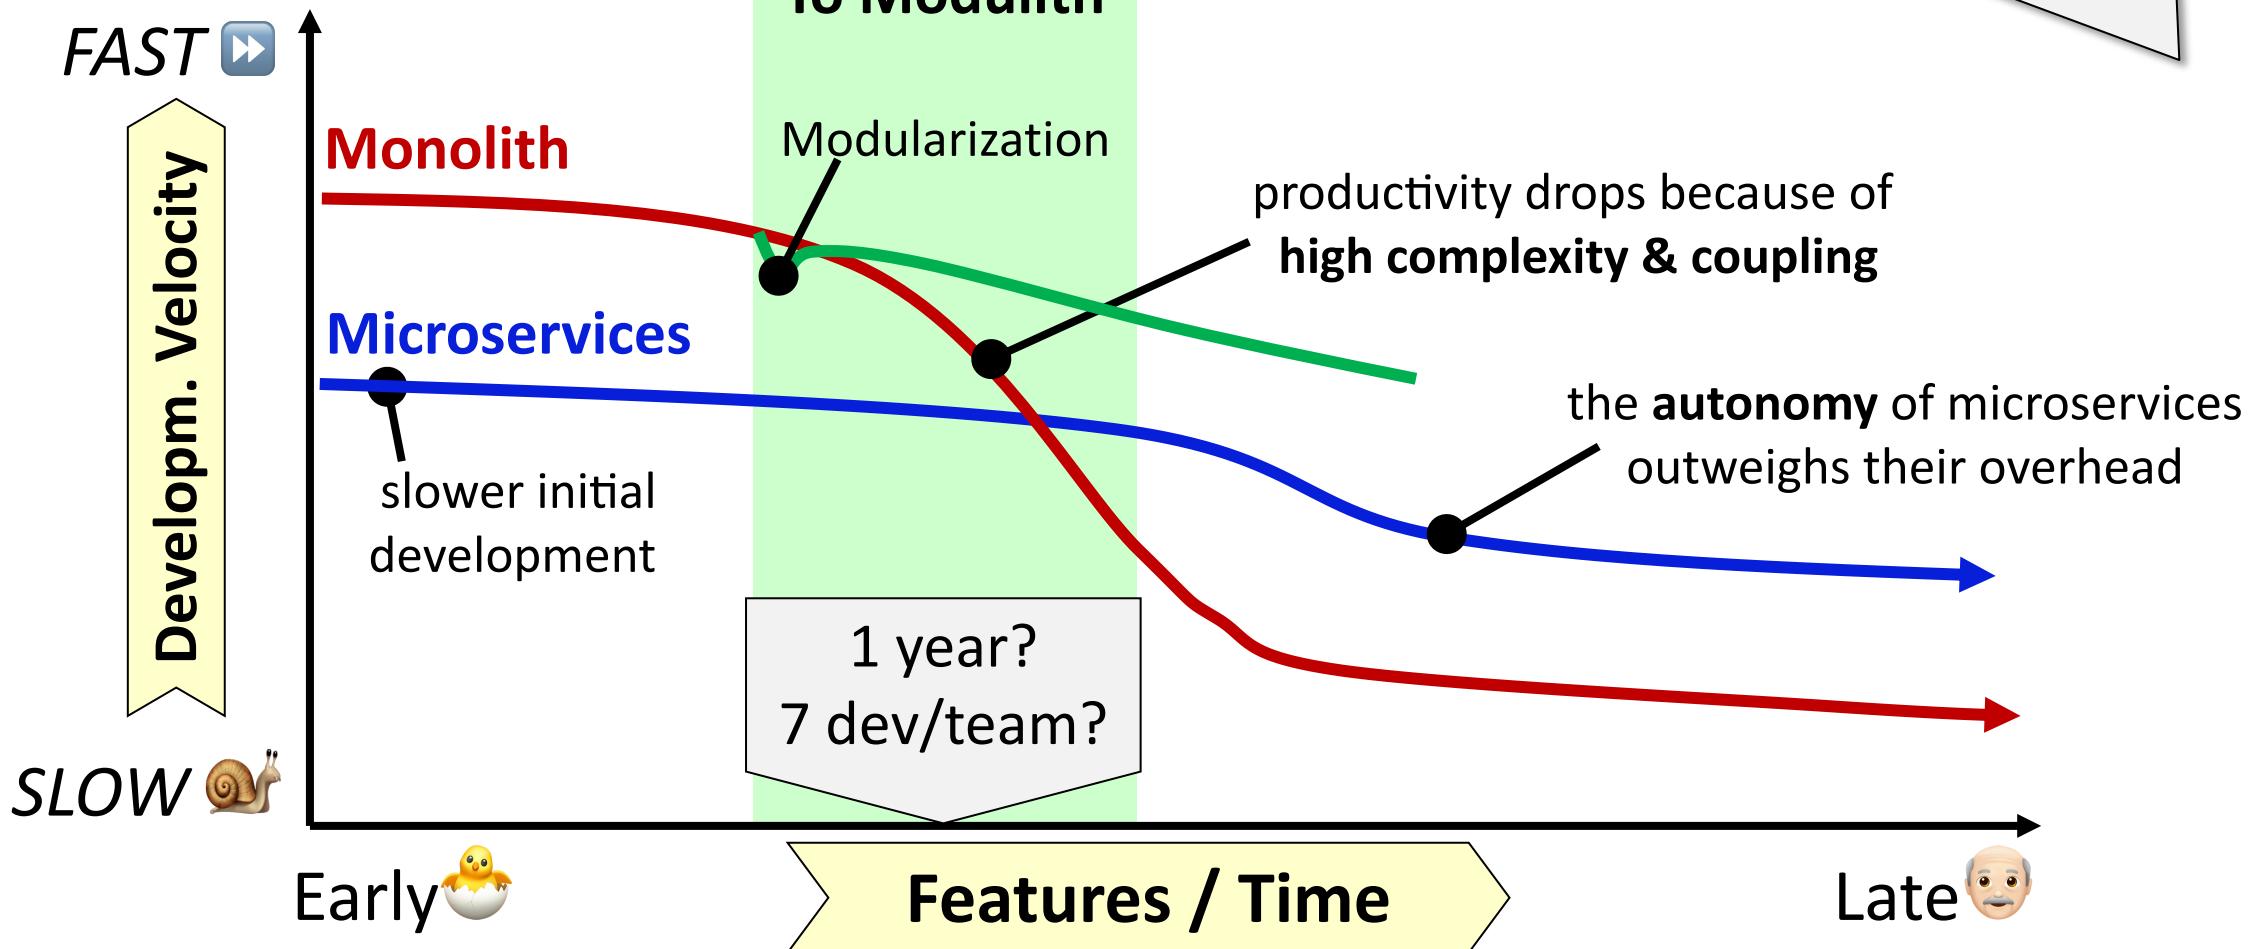
Evolutionary Patterns to Transform
Your Monolith



Sam Newman

Productivity

My system **will be complex**, so let's start with microservices!



You should **NOT** start a new project with microservices,
even if you're sure your application will be big enough
to make it worthwhile.

- Martin

Fowler

<https://martinfowler.com/bliki/MonolithFirst.html> (2015)
~~principal microservice evangelist~~

Instead: start simple and address
the system's **natural bottlenecks** as they occur
(functional, social, technical)



Meanwhile, in Reality: If we don't start with microservices, THEY (the bosses) won't give us the time we need to break it into microservices later = **WASTE**

Monolith **VS** Microservices

Can we have the best of both worlds?

Deployment architecture of a **Monolith**

Logical decoupling of **Microservices**

+ Ability to **rapidly extract** a Microservice out when necessary

Modular Monolith

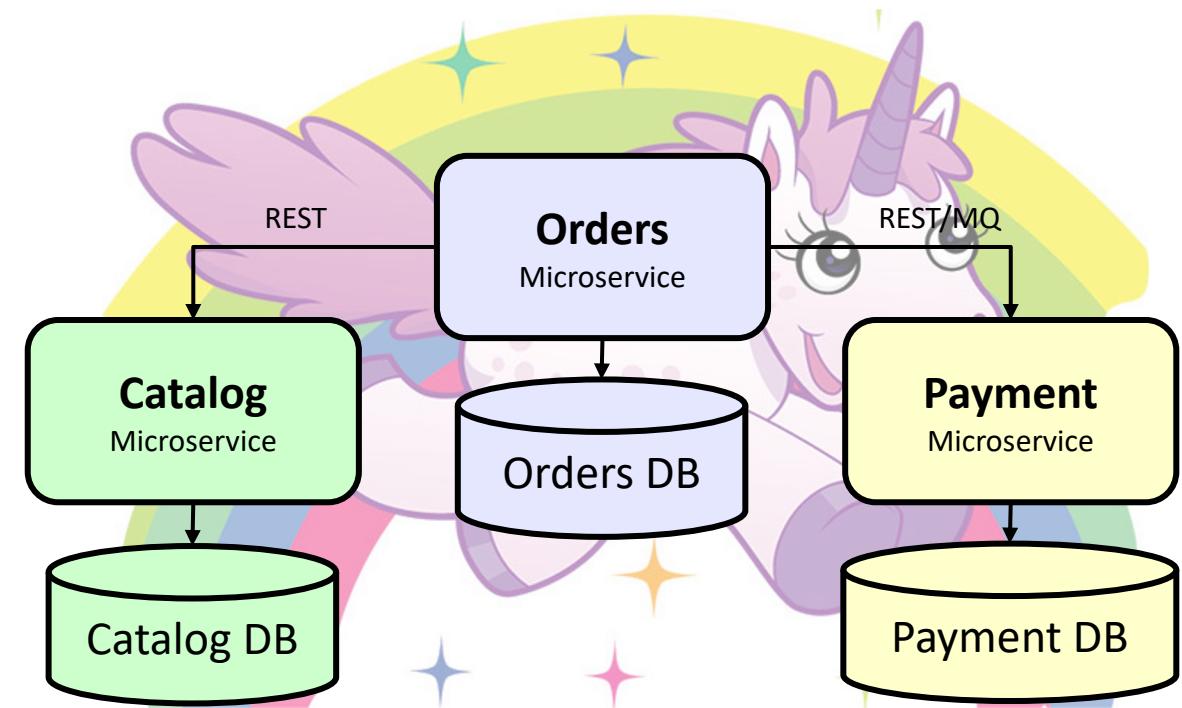
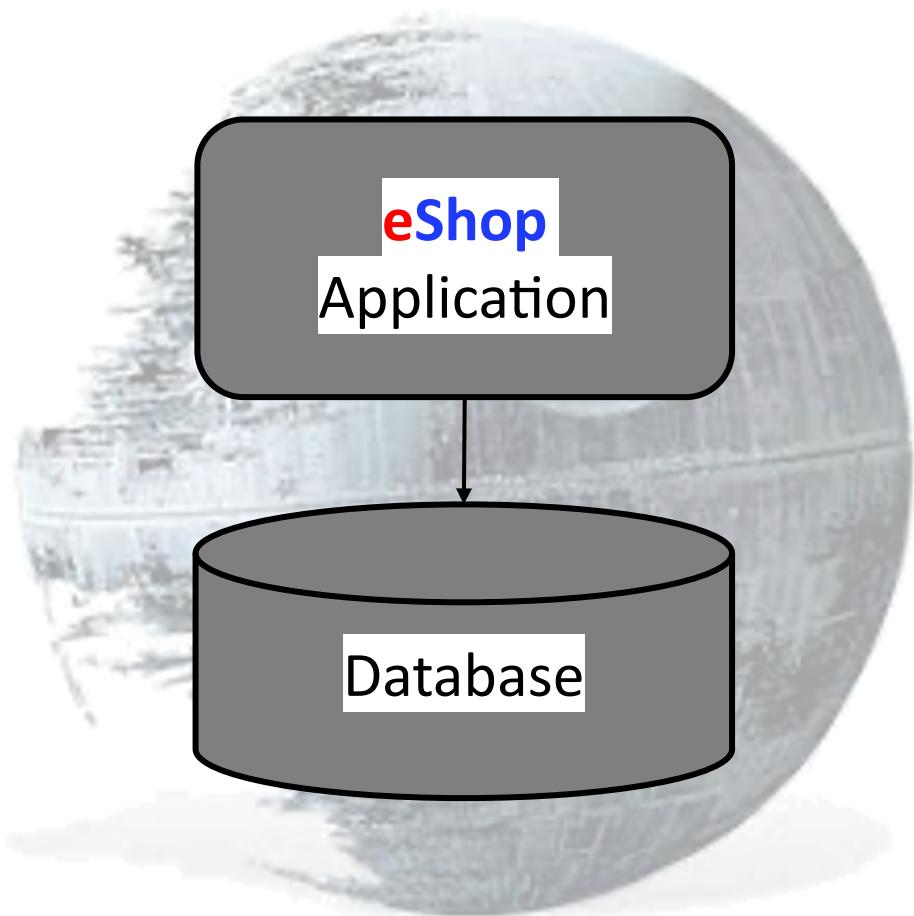


Modulith

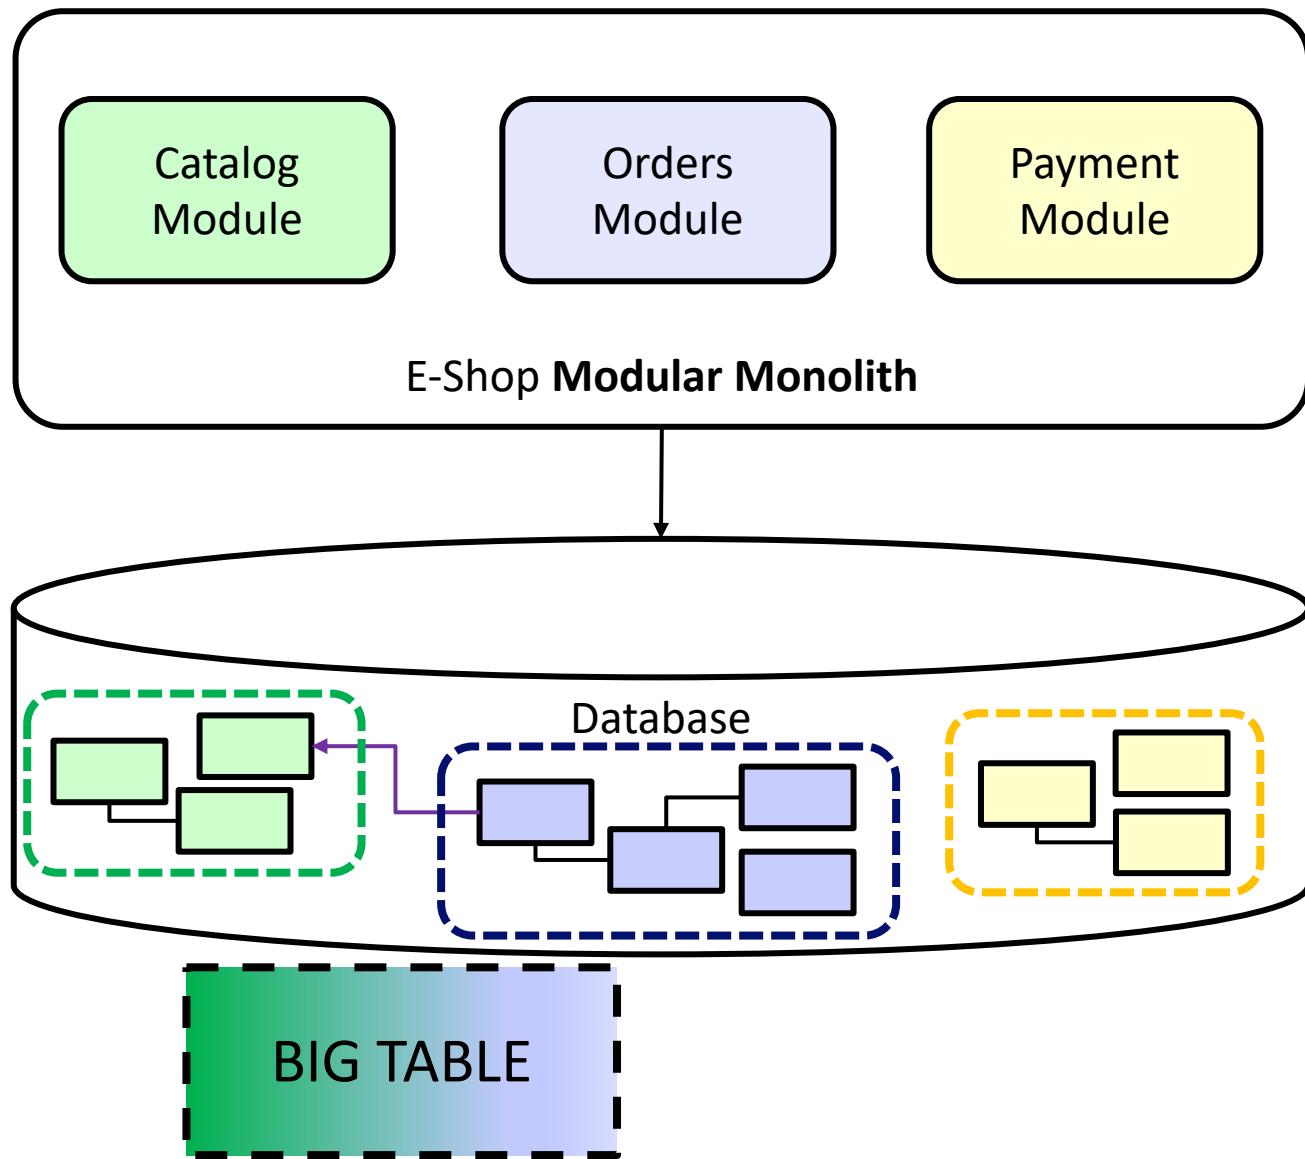
BBoM Monolith

VS

Microservices



Modular Monolith



What is a Module?

= a stand-alone logical application, having its own:

- **business features** to offer to users
- **private implementation**: domain model&logic, Rest Apis, infra
- **public API**:
 - **internal**: for other modules, via in-memory calls or events
 - **external**: for other systems, via REST, Rabbit, Kafka...
- **private persistence** in its own DB schema
- **micro-frontend**: screens & shared components (monorepo)

Life in a Modulith

independent teams working on decoupled modules

Example #1 (ehealth)

8 teams x 6-7pax {4-5 DEV+FE+PO+QA},
owning ~4 modules {with private DB schema}

+ 2 x Product Architects (Functional) + 3 x Core Team ("Platform")

CI build takes ≤15m (hard-core tuning involved)

Deploying in production 1/week

Example #2 (pharmacy)

5 teams x 3-4 DEV owning ~4 modules

CI = 30 minutes

CD = 2 weeks

Modulith Challenges

Finding correct module boundaries

Encapsulation of Modules

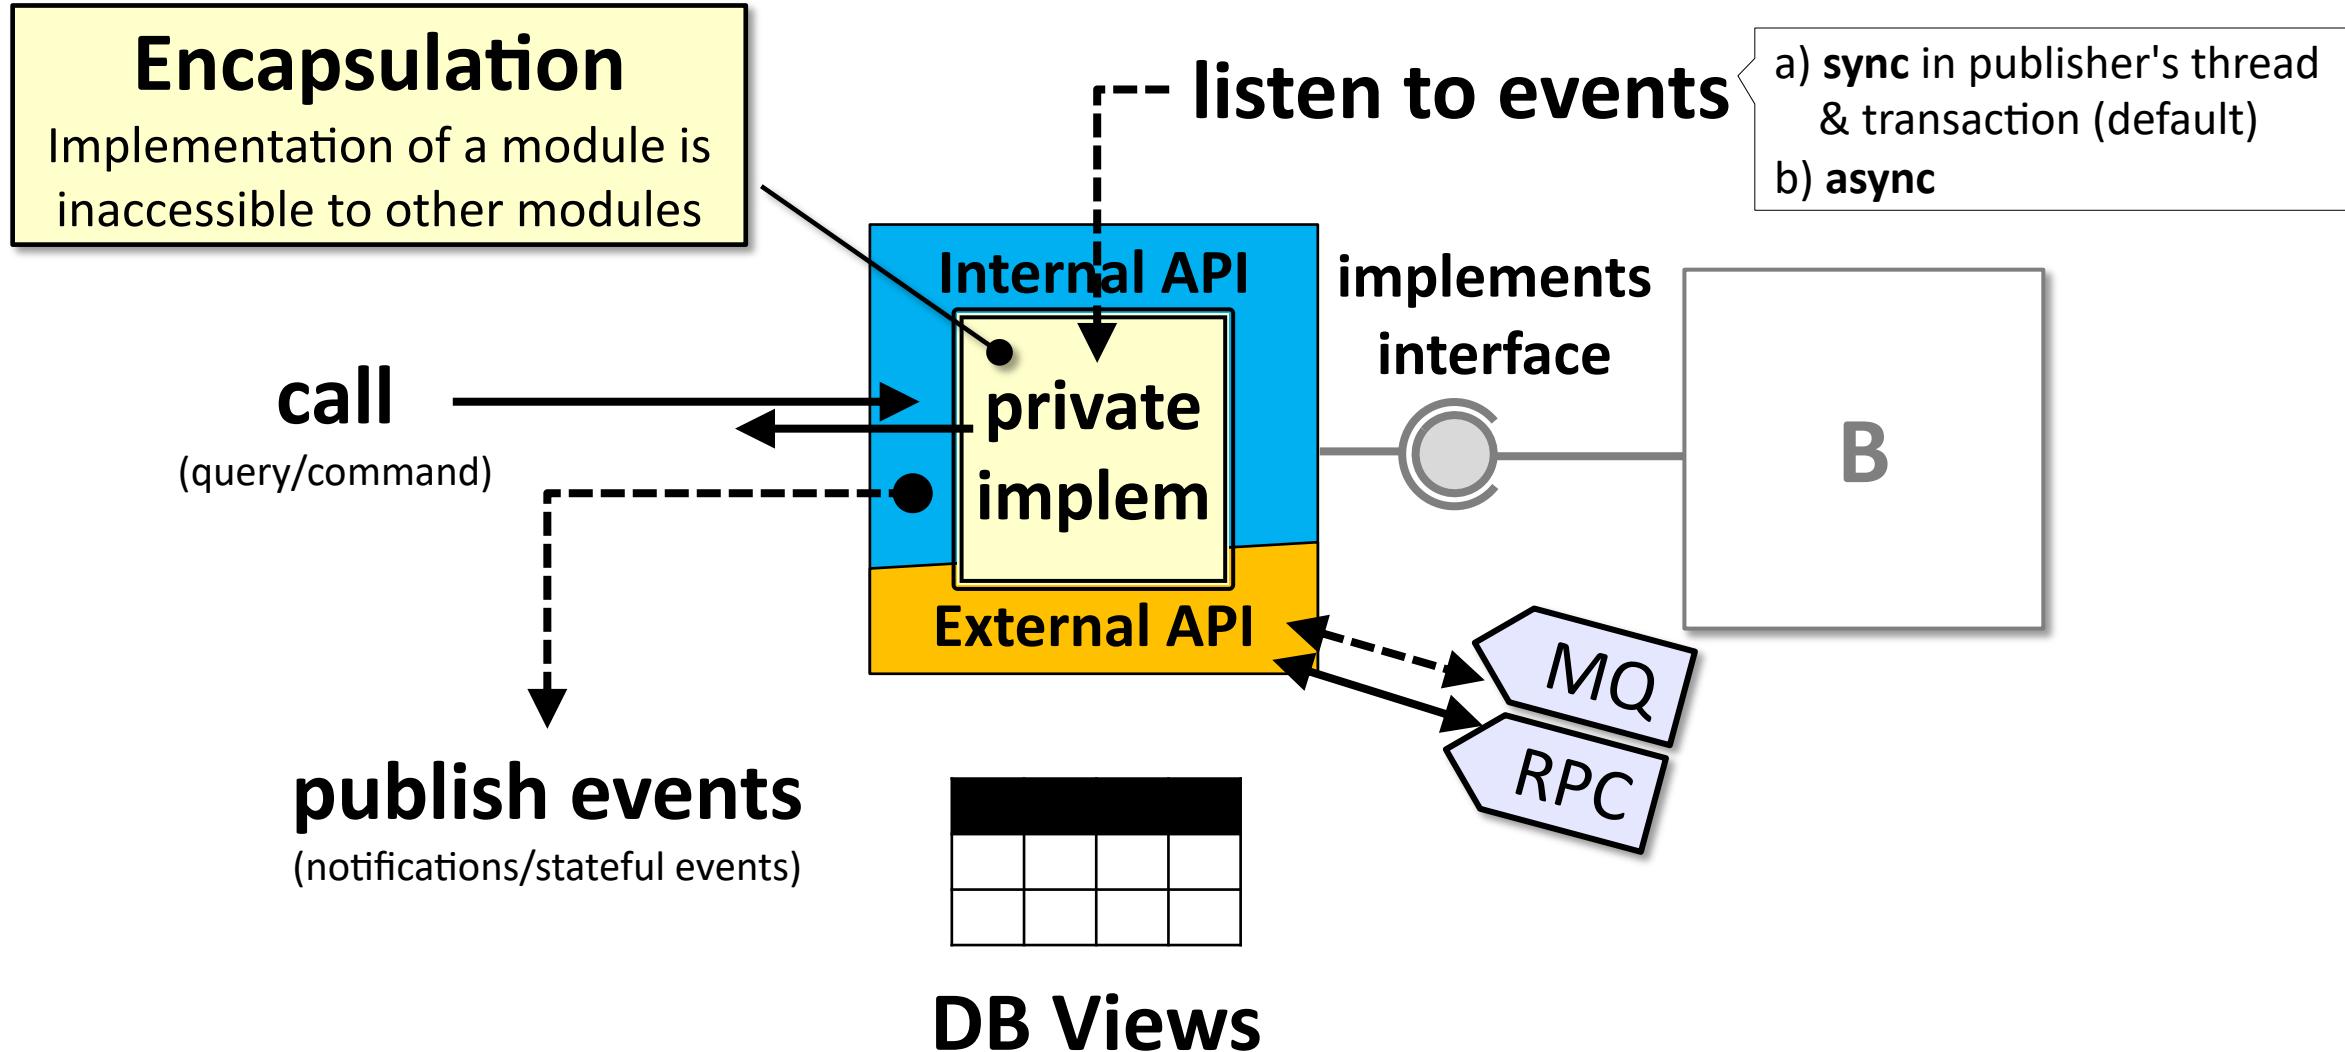
Data Decoupling

Testing

Module Public API

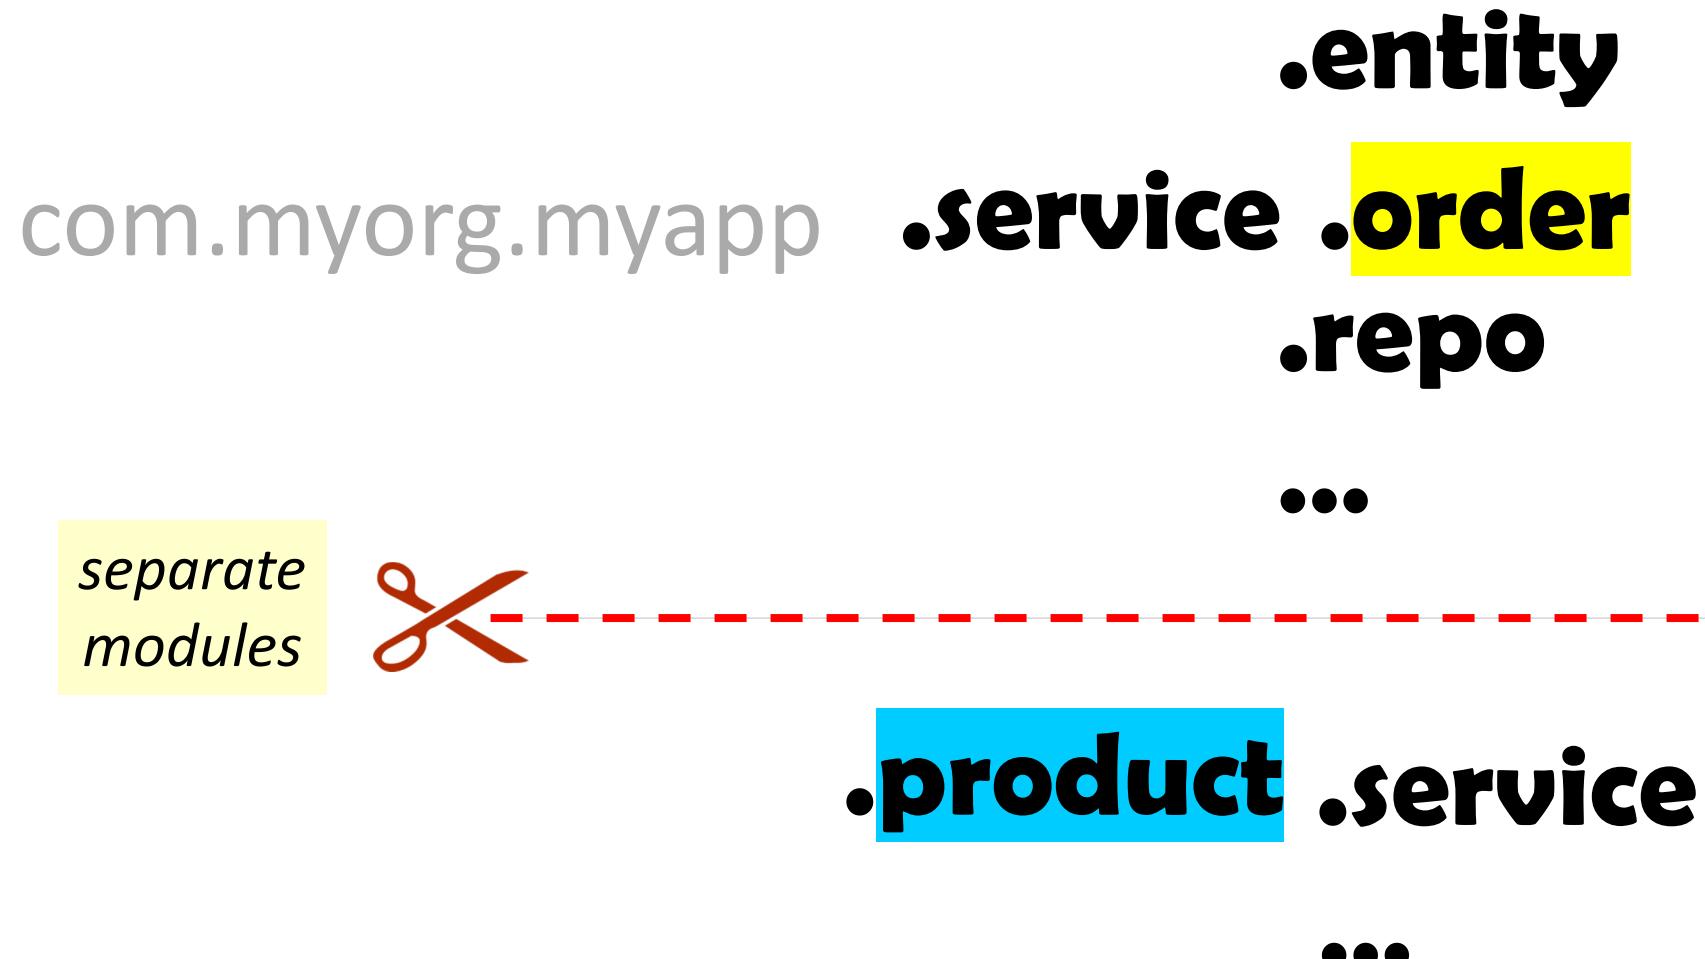
internal (for other modules)

external (for other systems)



The Screaming Architecture

Top-level packages reflect your domain structure



Enforce Boundaries

- Separate Build Units: eg Maven/Gradle modules
 - *customer-impl* depends on *customer-api* + the *order-api* it calls/listens
- ArchUnit www.archunit.org

```
@Test // unit test on CI
public void independentSubdomains() {
    slices().matching("..myapp.(*)..*")
        .should().notDependOnEachOther()
        .ignoreDependency(resideInAnyPackage("..shared.."))
        .check(new ClassFileImporter()).importPackages("com.myapp"));
}
```

- Spring Modulith spring.io/projects/spring-modulith

```
@Test
void verifyModularity() { // uses ArchUnit rules under the hood
    ApplicationModules.of(ModulithApp.class).verify();
}
```

Architecture is the art of deferring decisions

Why makes Modulith more Maintainable?

.... than a Big Ball of Mud?

1. Decoupled Logic via clear Module APIs

- Less accidental coupling
- Lower cognitive load per module

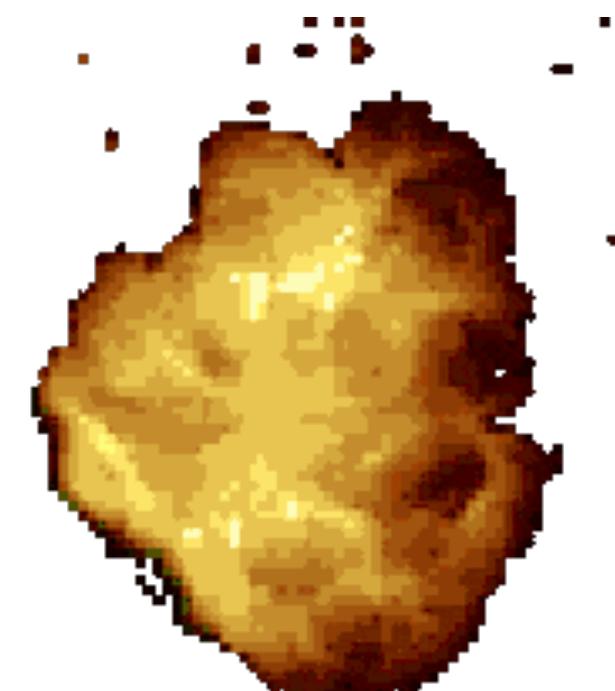
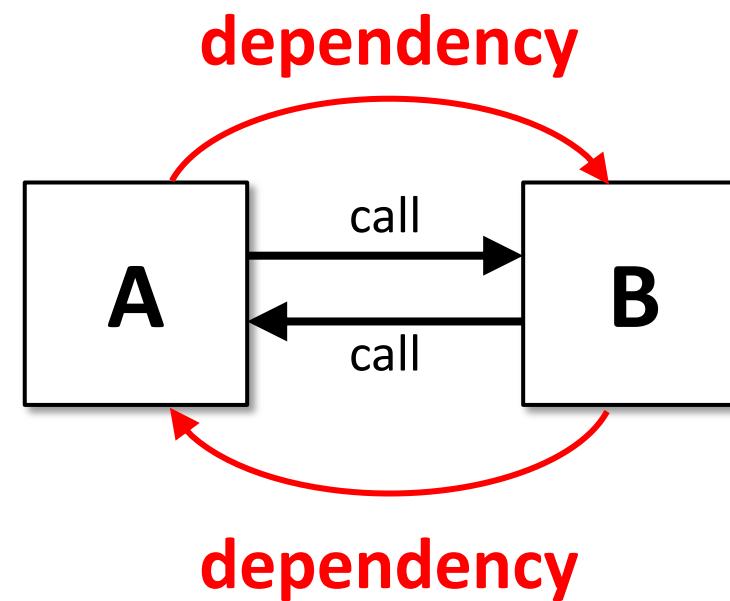
2. Smaller Domain Model can be more specialized

- catalog.Product vs order.Product

3. Separate Persistence in private schema/module

- More constraints in DB
- Control over tables: can keep them in sync with Domain remodeling

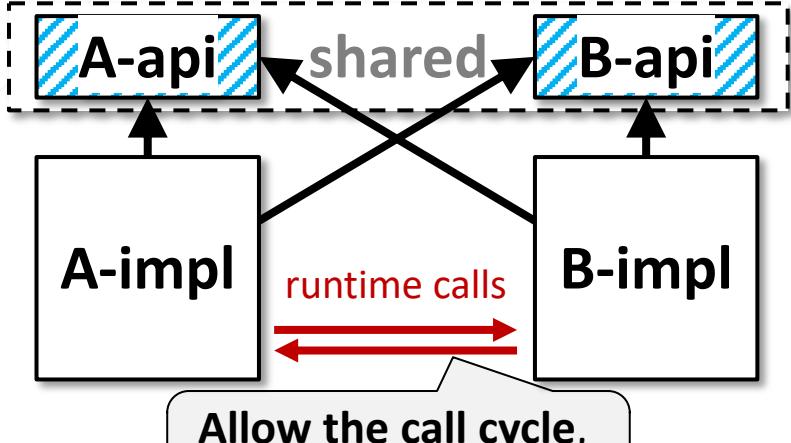
Circular dependencies between 2 modules



= Tight coupling

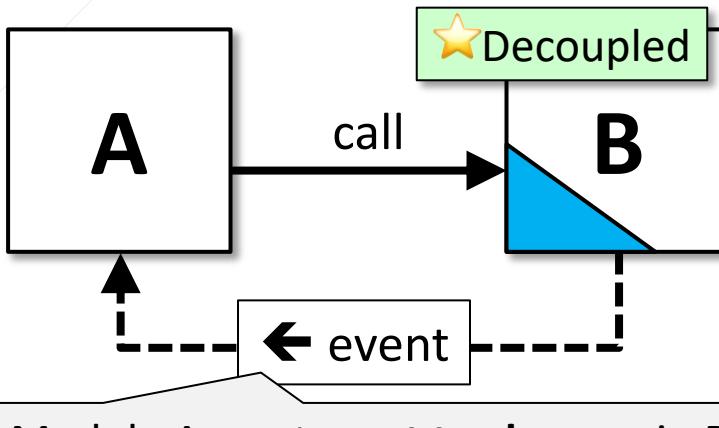
- Build fails for separate compilation units (eg maven/gradle)

Extract API as separate module



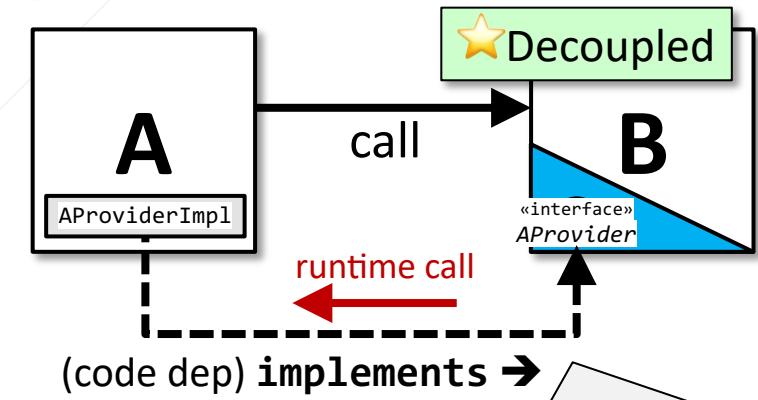
Allow the call cycle,
just fix the code cycle

Publish Events (in-mem)



Module A must react to changes in B

Dependency Inversion

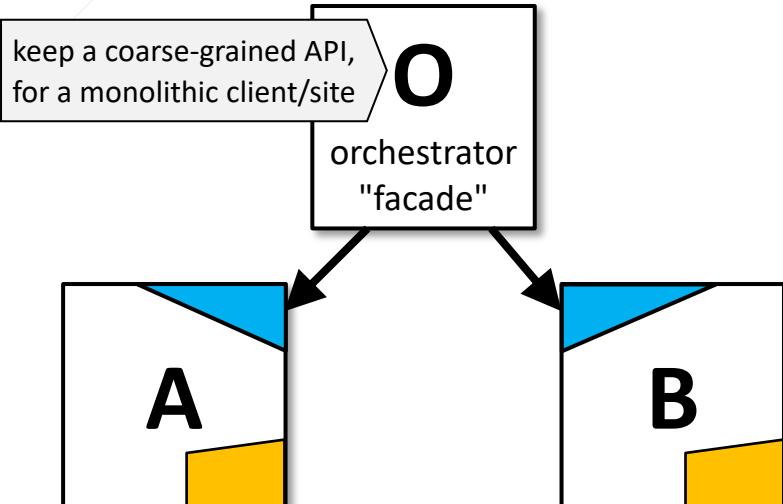


Using Dependency Inversion

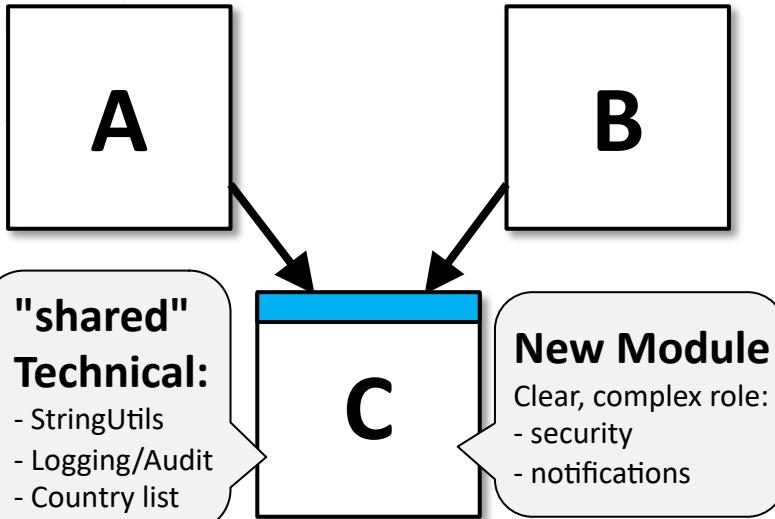
B can let A implement logic,
without depending on A

Strategically Dependencies between Modules

Pull Orchestration Up



Push Shared Down



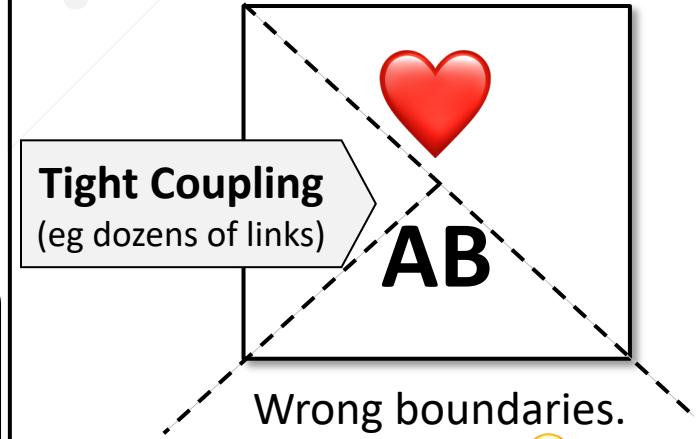
Tight Coupling
(eg dozens of links)



AB

Wrong boundaries.
Try others? 😐

Merge Modules



Module Events

- ⚠️ Events are part of the [Public Internal API](#) of the **publisher** module
- ⚠️ Publisher should expect nothing back (no result, no effect)
- 😊 Decoupled: publisher is unaware of listeners (potentially multiple)
- ⚠️ Order in which listeners run should NOT matter (chain events otherwise)
- ⚠️ Listeners run sequentially in the publisher's thread using current transaction (if any)
- 😊 Potentially @Async
- 😊 Can persist your events (see [@ApplicationModuleListener](#)) or send yourself a rabbit
- 😢 Events are harder to navigate than a direct method call → use events cross-module
- 🤔 Events can carry state to avoid a call back from listener ([Event-Carried State Transfer](#))

Microservices/Modules should have private tables

Allows you to keep in sync:
Problem \leftrightarrow Domain Model \leftrightarrow DB

Data Isolation Between Modules

KEEP OUT! CAUTION! KEEP OUT! CAUTION! KEEP OUT! CAUTION! KEEP OUT! CAUTION! KEEP OUT! CAUTION!

1. No Isolation (starting point): everyone freely **reads/writes** any table

⚠ Data Inconsistencies: a module could write valid data, but read garbage data.

2. Write Isolation: **one** module **writes** into a table, **any** other can **read**

🧠 Might require splitting a table (in separate schemas)

INVENTORY.PRODUCTS vs **CATALOG.PRODUCTS**, **CATALOG.PRODUCT_ATTRIBUTES**

⚠ Frozen Tables: an ALTER TABLE can break the readers

3. Exclusive Access 😇💖: a schema accessed by ONE single module

🧐 Trick: expose read-only VIEWS for others to JOIN : **INVENTORY STOCK_VIEW(id, items)** 😃
Strong Consistency

4. Consistency per Module - tables of two modules:

- (a) do NOT share any **Foreign Key**, and
- (b) are NOT updated in the same **Transaction**

Eventual Consistency

Defer until extraction as a
microservice is justified

Agile DB Evolution

DB incremental migration scripts, eg: liquibase/flyway/dbmaintain

- Never change a committed script, unless to fix a syntax error
- Unit tested on CI after every commit

```
v1_initial.sql:      create table product (id int8 not null, ...);  
v2_add_supplier.sql  create table supplier (id int8 not null, ...);  
v2_UNDO.sql          alter table product add constraint FK_PRODUCT_SUPPLIER ...  
                      alter table product add column supplier_id int8;
```

■ DB refactoring in baby-steps: (eg split an OLD table in 2 NEW)

- create 2 NEW tables; copy data from OLD to NEW ones
- have logic write in both OLD and NEW but read from OLD => prod
- 1 month later, if all looks good in NEW table in prod, make logic read from NEW => prod
- Trick 😊: Create a VIEW identical to OLD table to allow progressive code migration to 2 NEW
- 3 production month later, delete OLD table

■ UNDO scripts 😎 can quickly undo DB after a faulty deploy

Packages > Build Modules

(Maven/Gradle)

⚠ Only after module boundaries are clear AND the large team is about to split

■ Stronger separation

- Impossible to @Disable/add exclusions to the ArchUnit @Tests
- Impossible to have cycles
- Encapsulate implementation: (a) extract *x-api* module or (b) keep all apis in *shared* module
- Better Ctrl-Space suggestions (imagine: 2 classes named 'Product' in different modules)

■ Selective dependencies

- A module can decide: + *reactor*, + *poi*, + *jasperreports*, - *lombok*
- ! Libraries of all modules share the SAME VERSIONS

■ Partial packaging

- Client customizations: *invoicing-nhs.jar* is only included in the release for the NHS client
- Purchased features: *payment-exports.jar* (included only if paid)

Difficulties of "Living together" in 1 Process

■ **Fragile** vs crash / starvation

- Modules share the threads/DB connection pools → Monitoring +++

■ **Shared Transactions** (±)

- Exception/bug in listener module rolls back your transaction

■ **Classes with same name** in different packages

- Cumbersome code navigation + DI might fail ([Spring fix](#))

■ **Long CI build**

- Tendency for fine-grained **fragile** unit-tests → Module-tests!

■ **Single OpenAPI**

- Unique versioning? Type name collisions?



Steps to Extract a Module as a Microservice

■ Separate persistence

- No cross-schema SELECT. 🤔 Think...
- Drop all **Foreign Keys** to/from other modules' schemas 🤔
- Stop **sharing transactions** with other modules 🤔

■ Separate threads

- Turn method calls into **REST** localhost:8080 calls 🤔 (propagate metadata)
- Handle **events asynchronously**: @Async/@ApplicationModuleListener
- Send events via **external messaging**: Rabbit, Kafka.. 🤔

■ Move module to a separate Git/CI-CD/.pod...

- Extract internal API into a separate -client library

DON'T DO THIS
UNLESS
EXTRACTING THE
MODULE AS
MICROSERVICE
IS PLANNED IN
NEAR-FUTURE



How small should microservices/modules be?



Smaller!



Tiny!
PTSD after
BBoM?



Smaller pieces,
more coupling

The Fallacy of Nano-Services

The **microservices honeymoon is over.**

Uber is refactoring **thousands of microservices into a more manageable solution;**

Kelsey Hightower is predicting **monoliths are the future;**

Sam Newman is declaring that **microservices** should never be the default choice,
but rather a **last resort.**

Take-Away

Modular Monolith is

- **The safest way to transition to Microservices**
 - Cheap to experiment with boundaries
 - Allows to progressively enforce them
- **A valid alternative to microservices**
 - For highly complex domains with decent NFRs

Target a Modulith instead of Microservices?

- **YES✓**: A **greenfield** implementation
 - eg a new peer-to-peer lending platform in the early stages of growing into a full suite of financial services.
- **YES✓**: A system with a **low to moderate scale**
 - eg a movie ticket booking platform with several thousand users and handles a few million transactions per week.
- **YES✓**: **Non-complex business software platforms**
 - eg a notes/documents syncing and management platform for consumer apps.
- **YES✓**: **Oversized BBoM legacy** applications before breaking into independent services
 - eg an existing large-scale banking platform monolith

- **NO✗**: When parts of the platform are managed by multiple independent teams
 - Especially if each team governs its choice of underlying **technology** and runtimes;
- **NO✗**: If a **service chassis** can significantly reduce entry cost of microservices
 - New, large-scale platforms built in mature environments
- **NO✗**: **High scale and volume** of business + good understanding of architecture

https://www.ufried.com/blog/microservices_fallacy_9_moduliths/

<http://blogs.newardassociates.com/blog/2023/you-want-modules-not-microservices.html>

<https://www.techtarget.com/searchapparchitecture/tip/Understanding-the-modular-monolith-and-its-ideal-use-cases>

Testing a Modulith

End-to-end Tests (all modules)
are easier in a monolith

Module-scoped Tests ❤️

Integration test a Module in isolation

- mock public APIs of other modules
- mock the event bus

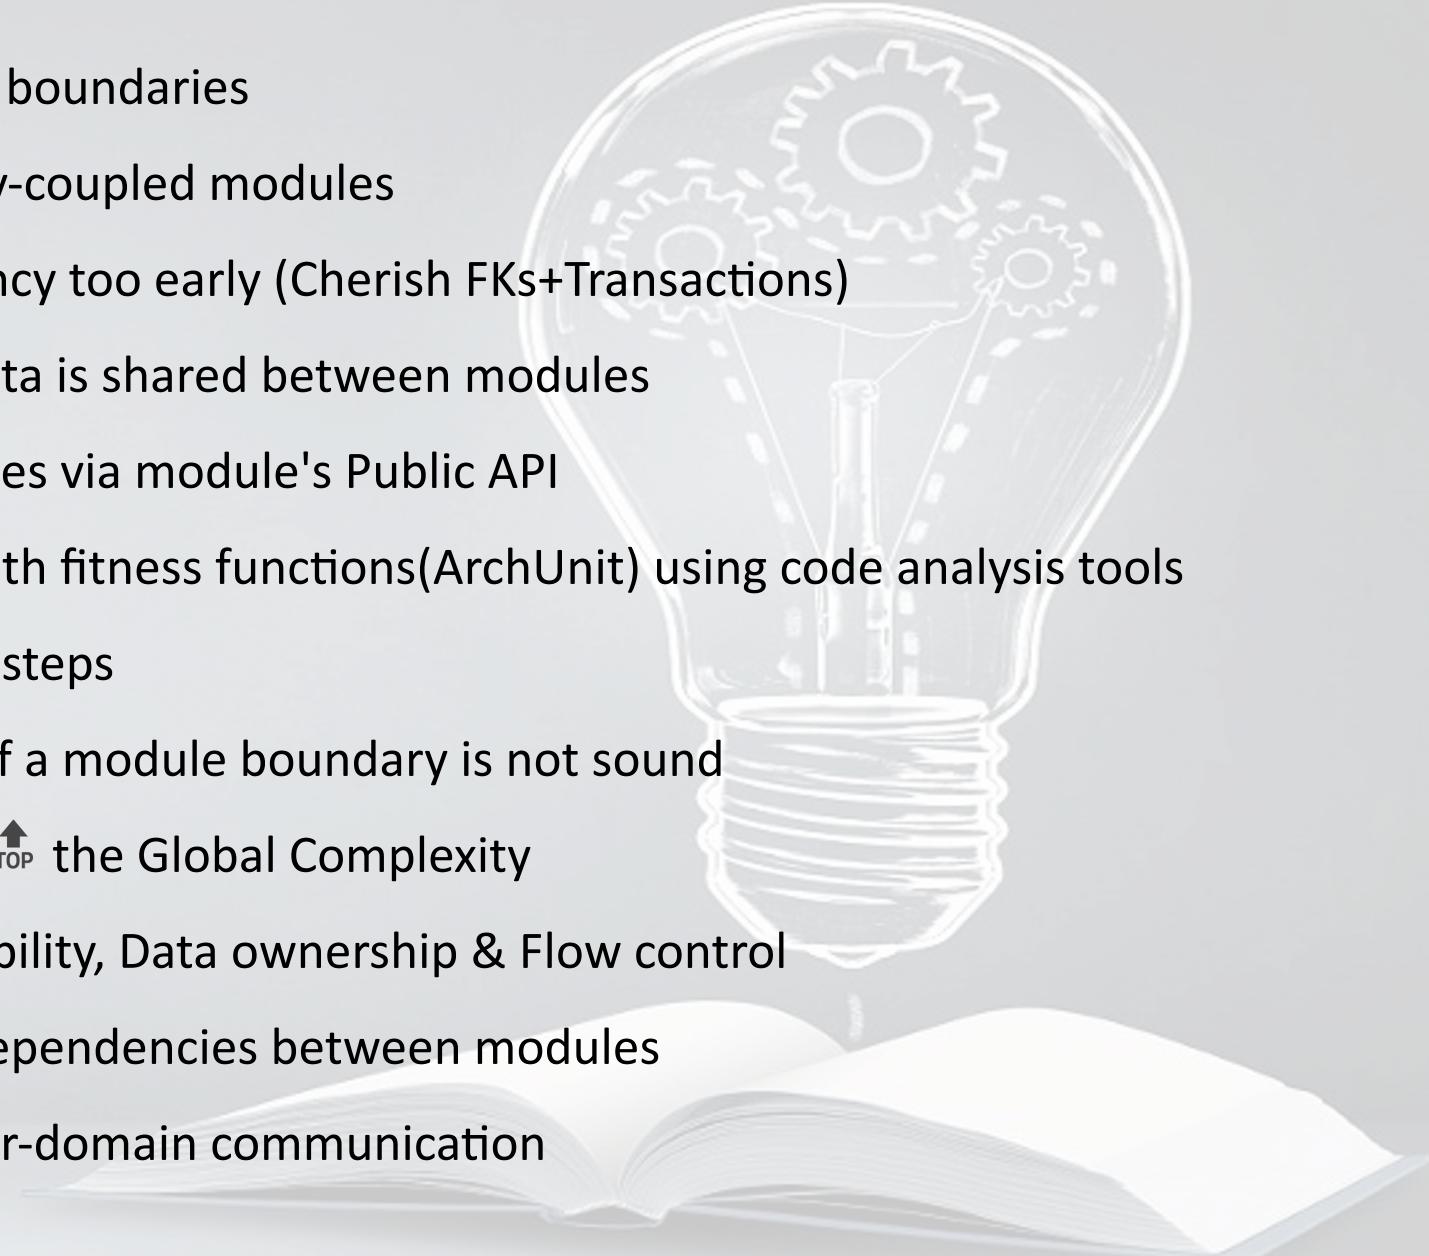
➔ See [Spring Modulith](#) for inspiration

Separate DB instances for module/micros?

- **(Yes) stricter control:** vs separate schema 
- **(Yes) fault tolerance:** DB crash should not impact other services
- **(Yes) scalability:** separate DB connections / service
- **(Yes) nosql is more suitable, eg:**
 - Redis, Hazelcast (in-memory, fast reads)
 - Mongo (document db)
 - Cassandra (heavy writes)
 - Neo4j (Graph DB), -specific
- **(No) hurts performance / too complex (legacy)**
 - Triggers or PL/SQL → pull logic out of DB
 - Views → move to eg. ElasticSearch
- **(No) strict consistency is IMPERATIVE for business (with arguments)**

Modulith – Lessons

- Carefully refine module boundaries
- Consider merging highly-coupled modules
- Don't sacrifice consistency too early (Cherish FKs+Transactions)
- Carefully design how data is shared between modules
- Don't leak internal classes via module's Public API
- Monitor architecture with fitness functions(ArchUnit) using code analysis tools
- Segregate code in baby-steps
- Be ready to stop/undo if a module boundary is not sound
- Too small modules will \uparrow_{TOP} the Global Complexity
- Focus on Business capability, Data ownership & Flow control
- Carefully refine cyclic dependencies between modules
- Consider Events for inter-domain communication

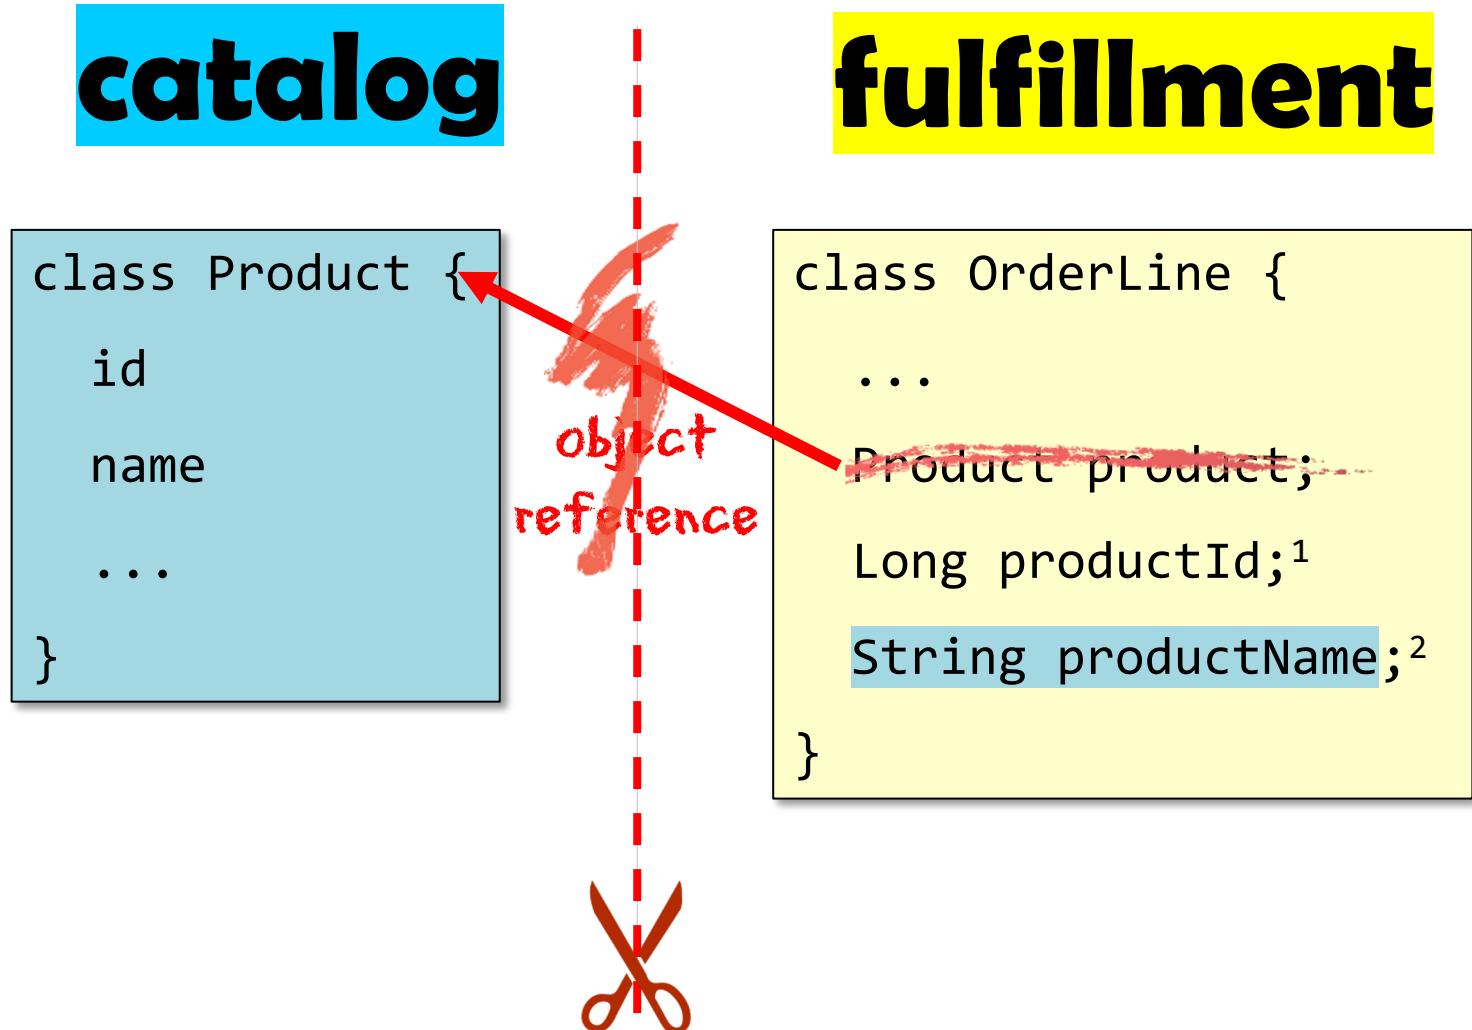


Towards Modules



in baby-steps

Decouple Domain Entities



Remove any object link to the Domain Model of another module:

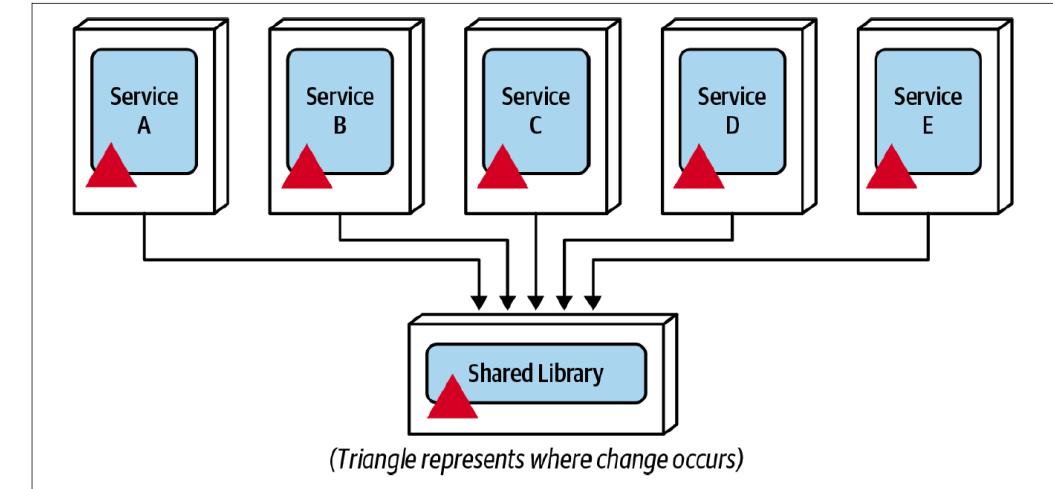
- 1) **Keep only the ID** ✓
👉 Preserve the FK (in a Modulith)
+ then later get from **catalog** on-demand
- 2) **Snapshot *immutable* data**
+ according to domain rules
- 3) **Keep *mutable* data in sync** 😰
+ replicate changes (eg via events)
⚠ Eventual Consistency (iff distributed)

Forces against extracting a Module out

- **Tight Coupling** by chained sync calls, especially bi-directional
 - Chained REST call would hurt ~~Availability~~ and ~~Performance~~

- **Shared Domain Model**

- BAD: Services depend on a 'commons' library holding 40% of business logic&data
 - You must change, test, deploy together
=> ~~Independent Deployability~~



- **Data Consistency** is critical between 2 Modules

- Many ACID flows and the complexity of sagas outweighs the benefits

Shared Library / Microservice Chassis

- **Problem:** implement common tasks in services without copy-paste
- **Solution:** create and manage a **shared library** helping with technical, non-domain specific tasks.
- **Debates** 🤔
 - Specialized libs ('-commons-security-v5') vs a single large one ('commons-v32')
 - Should clients use version=LATEST? Dependabot auto-upgrade?
 - What's the deprecation policy?
 - Who is paid to maintain the lib? Who's its GateKeeper?
- **Why turn a lib into a service?**
 - a) better support high rate of change
 - b) reuse across languages
 - c) clearer boundary and ownership

Finding Boundaries

between Modules / Microservices



= one of the hardest architecture challenges

<https://opencredo.com/blogs/identify-service-boundary-heuristics/>

<https://github.com/TeamTopologies/Independent-Service-Heuristics>

<https://www.dddheuristics.com/>

<https://www.michaelnygard.com/blog/2018/01/services-by-lifecycle/>

Boundary Types

- **Business:** subparts of the problem
 - eg: shipping, invoicing , ...
 - We speak the same **Ubiquitous Language** within the **Bounded Context** of our sub-problem
- **Humans (teams)**
 - Align **Dev Team** with **Business Team** (= **Conway's Law**)
- **Technical (codebase)**
 - Dev Team should have ownership on a manageable code part
 - ie. break the legacy Big Ball of Mud in Modules / Microsevices?
- You will get them wrong first → **PREPARE TO ITERATE** 

Group by Data (Entity Service)

= central data concepts

order

Order & friends

invoice

Invoice & friends

product

Product & friends

Group by Business Capability

= value to the users

(Feature Service)

fulfillment

place-
return-
workflow
history

invoicing

issue invoice
VAT discount
storno
taxes

catalog

search
filter
display
compare

Identify more
business capabilities
of an e-shop



Group by Business Capability

Example: product price

💡 Think in terms of attributes, not classes

```
class Customer {  
    firstName  
    lastName  
    status  
    ...  
}
```

```
class Product {  
    name  
    description  
    price  
    ...  
}
```



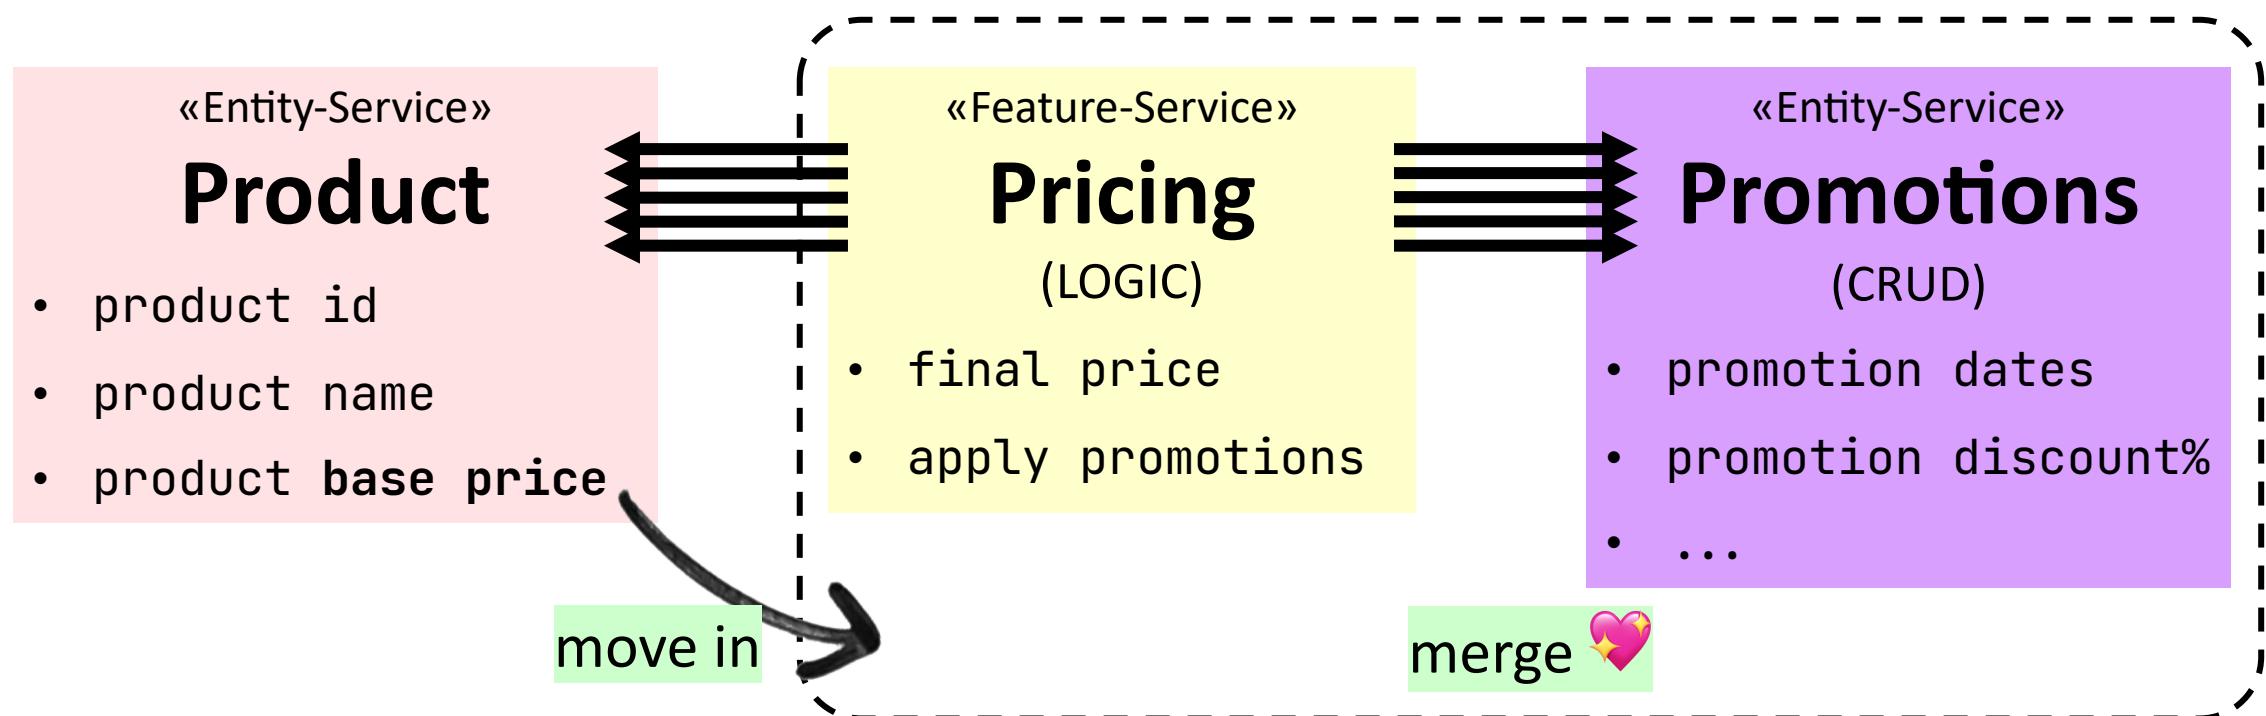
«Entity-Service»
Customer
• customer first name
• customer last name

«Feature-Service»
Pricing
• customer status
• product price
• promotions 😱😱😱

«Entity-Service»
Catalog
• product name
• product description

Group by Feature

Example: Promotions



Avoid Modules that are:

- Just **Behavior**

Example: calculation, validation

Except if they solve a very specialized domain problem.

- Just **CRUD Data**

Example: Country, Tenant, Clinic

→ Modulith: keep this data in a 'shared' module

→ Micros: inject static data as config

**A Module/Microservice is the
Technical Authority for a Business Capability**

